

# Data structuring, part 1

**The Pandas way**

*Andreas Bjerre-Nielsen*

# Recap

*Which Python containers have learned about so far?*

- ...
- ...

*Which containers can we turn into a `numpy` array?*

- ...

# Agenda

1. motivation
2. the numpy module
3. overview of the pandas module:
4. the pandas series
  - working with series and numeric procedures
  - boolean series
5. more tools:
  - inspecting and selecting observations
  - modifying DataFrames
  - dataframe IO

**Why we structure data**

# Motivation

*Why do we want to learn data structuring?*

- Data never comes in the form of our model. We need to 'wrangle' our data.

*Can our machine learning models not do this for us?*

- Not yet :). The current version needs **tidy** data. What is tidy?

One row per observation.

## Loading the software

```
In [33]: import numpy as np  
import pandas as pd
```

**Numpy module**

# Numpy overview

What is the numpy (<http://www.numpy.org/>) module?

numpy is a Python framework similar to matlab

- fast and versatile for manipulating arrays
- linear algebra tools available
- many machine learning and statistics libraries are built with numpy



# Numpy arrays

*What is a numpy array?*

An n-dimensional array:

- 1-d: vector;
- 2-d: matrix;
- 3-d: tensor.

We see use cases later.

# **Pandas overview**

# Pandas motivation

*Why use Pandas?*

1. simplicity - Pandas is built with Python's simplicity
2. flexible and powerful tools for manipulating data
3. speed - build on years of research about numeric computation
4. development - breathtaking speed of new tools coming

# Pandas data types

*How do we work with data in Pandas?*

- We use two fundamental data structures: `Series` and `DataFrame`.

# Pandas data frames (1)

*What is a DataFrame?*

- A matrix with labelled columns and rows (which are called indices). Example:

```
In [34]: df = pd.DataFrame([[1,2], [3, 4]],  
                           columns=['A', 'B'],  
                           index=['i', 'ii'])  
  
print(df)
```

	A	B
i	1	2
ii	3	4

- An object with powerful data tools.

## Pandas data frames (2)

*How are pandas dataframes built?*

Pandas dataframes can be thought of as numpy arrays with some additional stuff.

Most functions from numpy can be applied directly to Pandas. We can convert a DataFrame to a numpy matrix with values attribute.

```
In [35]: df.values
```

```
Out[35]: array([[1, 2],  
                [3, 4]], dtype=int64)
```

*To note:* In Python we can describe it as a list of lists of a dict of dicts.

# Pandas series

*What is a Series?*

- A vector/list with labels for each entry. Example:

```
In [36]: my_series = pd.Series([1,True,1/3,'a'])  
my_series
```

```
Out[36]: 0      1  
1      True  
2    0.333333  
3      a  
dtype: object
```

*What data structure does this remind us of?*

- A mix of Python list and dictionary (more info follows)

# Series vs DataFrames

*How are Series related to DataFrames?*

Every column is a series. Example: access as object method:

```
In [37]: print(df.A)
```

```
i      1  
ii     3  
Name: A, dtype: int64
```

Another option is access as key:

```
In [38]: print(df['B'])
```

```
i      2  
ii     4  
Name: B, dtype: int64
```

*To note:* The latter option more robust as variables named same as methods, e.g. count, cannot be accessed.





# Indices and column names

*Why don't we just use matrices?*

- inspection of data is quicker
- keep track of rows after deletion
- indices may contain fundamentally different data structures
  - e.g. time series, hierarchical groups
- facilitates complex operation:
  - merging datasets
  - split-apply-combine

# Working with pandas Series

# Generate Series (1)

Let's revisit our series

```
In [39]: my_series
```

```
Out[39]: 0          1
          1      True
          2    0.333333
          3          a
          dtype: object
```

Components in series

- index: label for each observation
- values: observation data
- dtype: the format of the series - object means any data type is allowed
  - note: the object dtype is SLOW!

## Generate Series (2)

*How do we set custom index?*

Example:

```
In [40]: num_data = range(0,3)
         indices = ['B','C','A']
         my_series2 = pd.Series(num_data, index=indices)
         my_series2
```

```
Out[40]: B      0
         C      1
         A      2
         dtype: int64
```

## Generation Series (3)

*Can a dictionary be converted to a series?*

Yes, we just put into the Series class constructor. Example:

```
In [41]: d = {'yesterday':0, 'today':1, 'tomorrow':3}

my_series3 = pd.Series(d)
my_series3
```

```
Out[41]: today          1
tomorrow         3
yesterday         0
dtype: int64
```

Note: Same is true for DataFrames which requires that each value in the dictionary is also a dictionary.

## Generation Series (4)

*Can we convert series to dictionaries?*

- Yes, in most cases.
- **WARNING!#@:** Series indices are NOT unique

```
In [42]: pd.Series(range(3),index=['A', 'A','A']).to_dict()
```

```
Out[42]: {'A': 2}
```

# The power of pandas

*How is the series different from a dict?*

- We will see that pandas Series have powerful methods and operations.
- It is both key and index based.



# Converting data types

The data type of a series can be converted with the **astype** method:

```
In [43]: my_series3.astype(np.float64) # np.str
```

```
Out[43]: today          1.0  
tomorrow         3.0  
yesterday        0.0  
dtype: float64
```

# **Numeric procedures**

# Numeric operations (1)

*How can we basic arithmetic operations with arrays, series and dataframes?*

Like Python data! An example:

```
In [44]: my_arr1 = np.array([2, 3, 2])  
         my_arr2 = my_arr1 ** 2  
         my_arr2
```

```
Out[44]: array([4, 9, 4], dtype=int32)
```

## Numeric operations (2)

*Are other numeric python operators the same??*

Yes `/`, `//`, `-`, `*`, `**` etc. behave as expected.

*Why is this useful?*

- vectorized operations are VERY fast;
- requires very little code.

## Numeric operations (3)

*Can we do the same with two vectors?*

- Yes, we can also do elementwise addition, multiplication, subtractions etc. of series. Example:

```
In [45]: my_arr1 + my_arr2
```

```
Out[45]: array([ 6, 12,  6])
```

# Numeric methods (1)

Pandas series has powerful numeric methods built-in, example:

```
In [46]: my_series2.median()
```

```
Out[46]: 1.0
```

Other useful methods include: **mean**, **quantile**, **min**, **max**, **var**, **describe**, **quantile** and many more.

```
In [47]: my_series2.describe()
```

```
Out[47]: count      3.0  
         mean       1.0  
         std        1.0  
         min        0.0  
         25%        0.5  
         50%        1.0  
         75%        1.5  
         max        2.0  
         dtype: float64
```

## Numeric methods (2)

An important method is `value_counts`. This counts number for each observation.

Example:

```
In [48]: my_series4 = pd.Series(my_arr2)
         my_series4.unique()
```

```
Out[48]: array([4, 9], dtype=int64)
```

```
In [49]: my_series4.value_counts()
```

```
Out[49]: 4    2
         9    1
         dtype: int64
```

What is observation in the `value_counts` output - index or data?

## Numeric methods (3)

*Are there other powerful numeric methods?*

Yes: examples include

- **unique, nunique**: the unique elements and the count of unique elements
- **cut, qcut**: partition series into bins
- **diff**: difference every two consecutive observations
- **cumsum**: cumulative sum
- **nlargest, nsmallest**: the n largest elements
- **idxmin, idxmax**: index which is minimal/maximal
- **corr**: correlation matrix

Check [series documentation \(https://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.html\)](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.html) for more information.



# Boolean Series

# Logical expression for Series

*Can we test an expression for all elements?*

Yes: `==`, `!=` work for a single object or Series with same indices. Example:

```
In [50]: my_series3 > 0
```

```
Out[50]: today          True
tomorrow          True
yesterday        False
dtype: bool
```

What datatype is returned?

## Logical expression in Series (2)

*Can we check if elements in a series equal some element in a container?*

Yes, the `isin` method. Example:

```
In [52]: my_rng = range(2)
          print(list(my_rng))
          # my_series3.isin(my_rng)
```

```
[0, 1]
```

# Power of boolean series (1)

*Can we combine boolean Series?*

Yes, we can use:

- the and operator, also known as &
- the or operator, also known as |

```
In [53]: # (my_series3 > 0) & (my_series3 == 1)
```

What datatype is returned?

## Power of boolean series (2)

*Why do we care for boolean series (and arrays)?*

Because we can use the to select rows based on their content.

```
In [54]: # my_series6 = pd.Series(data=[17, 18, 18],
#                                     index=['April 1', 'April 2', 'April 3'],
#                                     name='age')
# print(my_series6)
# my_series6[my_series6>17]
```

NOTE: Boolean selection is extremely useful for dataframes!!

**Inspecting and selecting observations**

# Viewing series and dataframes

*How can we view the contents in our dataset?*

- We can use `print` our dataset
- We can visualize patterns by plotting (from tomorrow)

# The head and tail

We select the *first* rows in a DataFrame or Series with the head method.

```
In [55]: n = 3 # number of observations  
my_series7 = pd.Series(np.random.normal(size=[100]))  
my_series7.head(n)
```

```
Out[55]: 0    0.128634  
1    1.129369  
2    1.153792  
dtype: float64
```

The `tail` method selects the last observations in a DataFrame.



## Row selection (1)

*How can we select certain rows in a Series when for given index **keys**?*

With the `loc` attribute. Example:

```
In [56]: my_loc = 'tomorrow'
          # my_loc = ['today', 'tomorrow']
          # my_series3.loc[my_loc]
```

## Row selection (2)

*How can we select certain rows in a Series when for given index **integers**?*

The `iloc` method selects rows for provided index integers.

```
In [57]: my_series3.iloc[1]  
# my_series7.iloc[10:13]
```

```
Out[57]: 3
```

## Row selection (3)

*Do our tools for viewing specific rows, i.e. `Loc`, `iLoc` work for DataFrames?*

- Yes, we can use both `loc` and `iloc`. As default they work the same.

```
In [60]: my_idx = ['i', 'ii', 'iii']  
my_cols = ['a', 'b']  
my_data = [[1, 2], [3, 4], [5, 6]]  
my_df = pd.DataFrame(my_data, columns=my_cols, index=my_idx)  
my_df.loc[['i']]
```

Out[60]:

	a	b
i	1	2

## Row selection (4)

*How are Loc, iLoc different for DataFrames?*

- For DataFrames we can also specify columns.

```
In [59]: idx_keep = ['i','ii']  
         cols_keep = ['a']  
         print(my_df.loc[idx_keep, cols_keep])
```

```
      a  b  
i      1  2  
ii     3  4  
iii    5  6  
      a  
i      1  
ii     3
```

# Columns selection

*How can we select columns in a DataFrame?*

- Option 1: using the `[]` and providing a list of columns.
- Option 2: using `loc` and setting row selection as `:`.

```
In [79]: # my_df[cols_keep]
```

## Selection quiz

*What does : do in iLoc or Loc?*

Select all rows (columns).

# Modifying DataFrames

## Chaging the index (1)

*How can we change the index of a DataFrame?*

We change set a DataFrame's index index using its method `set_index`. Example:

```
In [87]: # my_df.set_index('a')
```



## Chaging the index (2)

*Is our DataFrame changed? I.e. does it have a new index?*

We can use the keyword `inplace` which will replace the DataFrame:

```
In [65]: my_df_a = my_df.set_index('a')  
my_df_copy = my_df.copy()  
my_df_copy.set_index('a', inplace=True)  
my_df_copy.head(2)
```

Out[65]:

	<b>b</b>
<b>a</b>	
<b>1</b>	2
<b>3</b>	4

## Chaging the index (3)

Sometimes we wish to remove the index. This is done with the `reset_index` method:

```
In [66]: my_df.reset_index()  
my_df
```

Out[66]:

	a	b
0	1	2
1	3	4
2	5	6

By specifying the keyword `drop=True` we delete the index. Note `inplace` also works.

*To note:* Indices can have multiple levels, in this case `level` can be specified to delete a specific level.

# Sorting data

A DataFrame can be sorted with `sort_values`; this method takes one or more columns to sort by.

```
In [85]: my_df.sort_values(by='a', ascending=False)
```

Out[85]:

	a	b
iii	5	6
ii	3	4
i	1	2

*To note:* Many key word arguments are possible for `sort_values`, including `ascending` if for one or more valuable we want descending values. Sorting by index is possible with `sort_index`.

## **DataFrame IO: loading and storing**

# Reading DataFrames (1)

Download the file from url:

```
In [76]: url = 'https://api.statbank.dk/v1/data/FOLK1A/CSV?lang=en&Tid=*
```

## Reading DataFrames (2)

Now let's try opening it:

- As local file:

```
In [84]: abs_path = 'C:/Users/bvq720/Downloads/FOLK1A.csv' # absolute path  
# rel_path = 'FOLK1A.csv' # relative path  
  
df = pd.read_csv(abs_path) # open the file as dataframe
```

What are absolute and relative paths?

## Reading DataFrames (3)

Now let's try opening it online file:

```
In [81]: df = pd.read_csv(url, sep=';')  
df.head(3)
```

Out[81]:

	TID	INDHOLD
0	2008Q1	5475791
1	2008Q2	5482266
2	2008Q3	5489022

## Reading other data types

Other pandas readers include: excel, sql, sas, stata and many more.

*To note:* an incredibly fast and useful module for reading and writing data is feather (<https://github.com/wesm/feather>).



## Storing data

Data can be stored in a particular format with `to_(FORMAT)` where (FORMAT) is the file type such as csv. Let's try with `to_csv`:

```
In [14]: df.to_csv('DST_people_count.csv', index=False)
```

Should we always set `index=False`. Yes, unless time series!!! Otherwise the index will be exported too!

# The end

[Return to agenda](#)